

Machine Learning and Data Mining

Linear regression

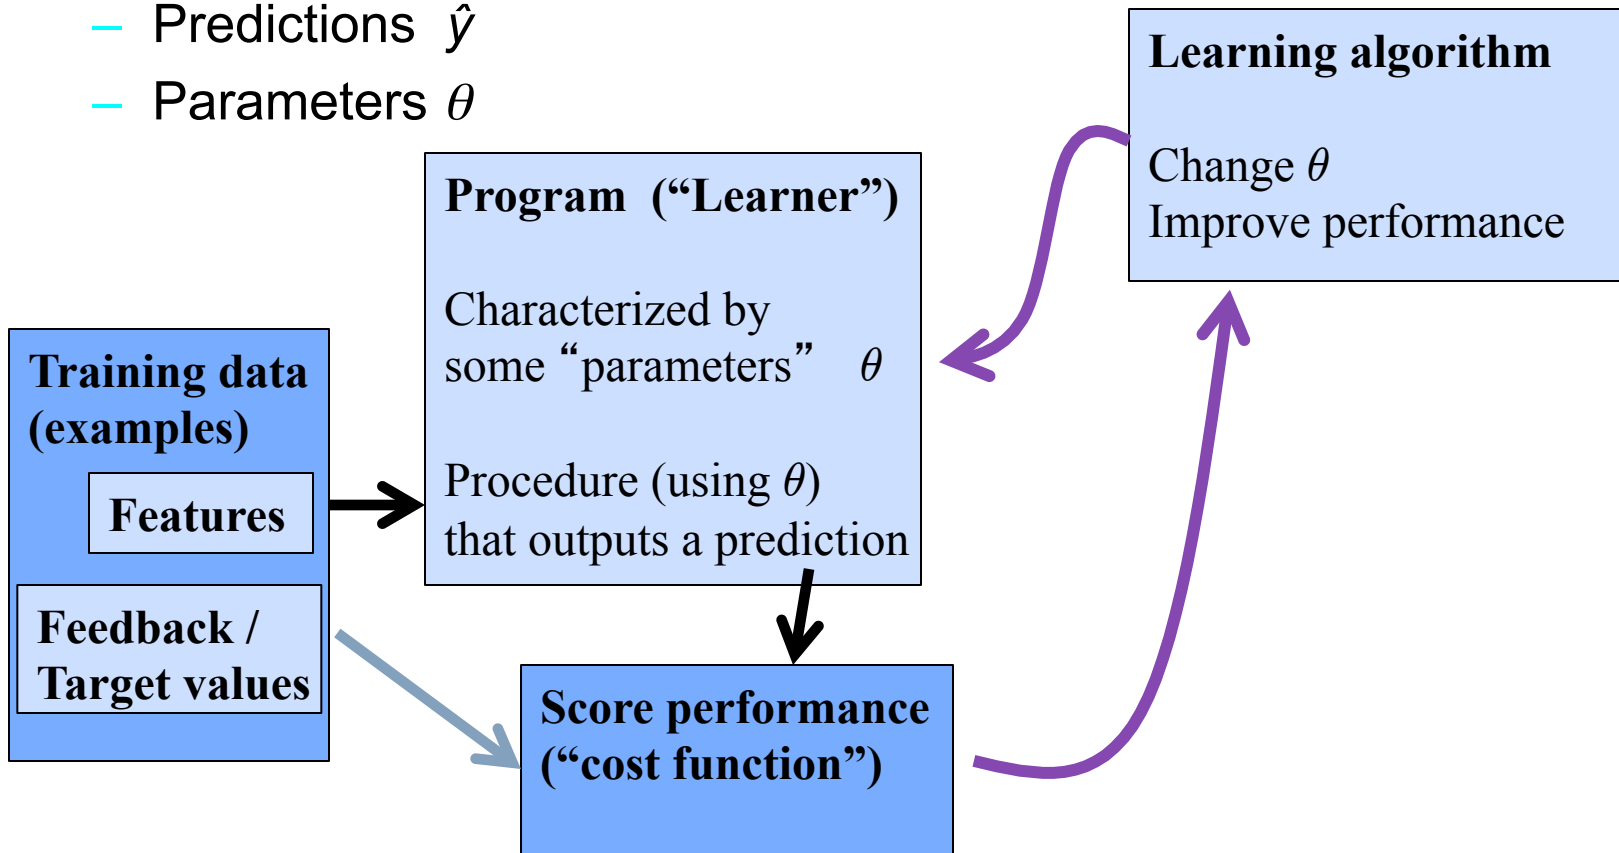
Prof. Alexander Ihler
Fall 2012




Supervised learning

- Notation

- Features x
- Targets y
- Predictions \hat{y}
- Parameters θ



Linear regression



“Predictor”:

Evaluate line:

$$r = \theta_0 + \theta_1 x_1$$

return r

- Define form of function $f(x)$ explicitly
- Find a good $f(x)$ within that family

Notation


$$\hat{y}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots$$

Define “feature” $x_0 = 1$ (constant)

Then

$$\hat{y}(x) = \underline{\theta} \underline{x}^T$$
$$\underline{\theta} = [\theta_0, \dots, \theta_n]$$
$$\underline{x} = [1, x_1, \dots, x_n]$$

Measuring error



Sum of squared error

- How can we quantify the error?

$$\begin{aligned}\text{SSE, } J(\underline{\theta}) &= \frac{1}{2} \sum_j (y^{(j)} - \hat{y}(x^{(j)}))^2 \\ &= \frac{1}{2} \sum_j (y - \underline{\theta} \cdot \underline{x}^T)^2\end{aligned}$$

- Could choose something else, of course...
 - Computationally convenient (more later)
 - Measures the variance of the residuals
 - Corresponds to Gaussian model of “noise”

$$\mathcal{N}(y ; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{1}{2\sigma^2} (y - \mu)^2 \right\}$$

SSE cost function

$$\begin{aligned}\text{SSE, } J(\underline{\theta}) &= \frac{1}{2} \sum_j (y^{(j)} - \hat{y}(x^{(j)}))^2 \\ &= \frac{1}{2} \sum_j (y - \underline{\theta} \cdot \underline{x}^T)^2\end{aligned}$$

- Rewrite using matrix form

$$\begin{aligned}\underline{\theta} &= [\theta_0, \dots, \theta_n] \\ \underline{y} &= [y^{(1)} \dots, y^{(m)}] \\ \underline{X} &= \begin{bmatrix} x_0^{(1)} & \dots & x_n^{(1)} \\ \vdots & \ddots & \vdots \\ x_0^{(m)} & \dots & x_n^{(m)} \end{bmatrix}\end{aligned}$$

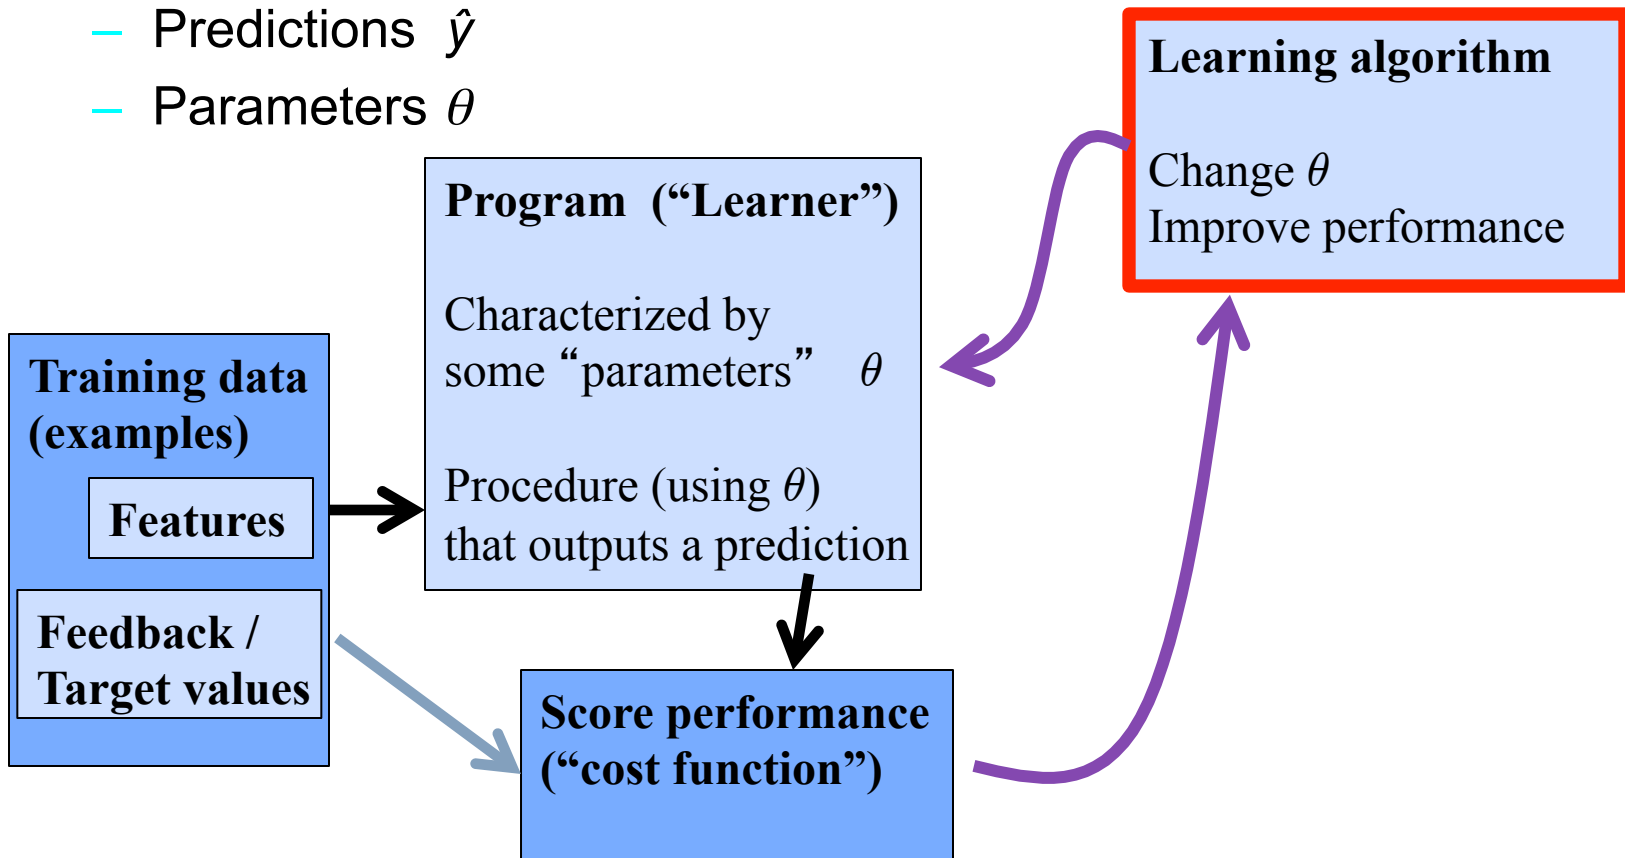
$$J(\underline{\theta}) = \frac{1}{2} (\underline{y} - \underline{\theta} \underline{X}^T) \cdot (\underline{y} - \underline{\theta} \underline{X}^T)^T$$

(Matlab) `>> e = y - th*x'; J = .5*e*e';`





Supervised learning

- Notation

- Features x
- Targets y
- Predictions \hat{y}
- Parameters θ




Visualizing the cost function




Finding good parameters

- Want to find parameters which minimize our error...
- Think of a cost “surface”: error residual for that θ ...




Gradient descent



- How to change θ to improve $J(\theta)$?
- Choose a direction in which $J(\theta)$ is decreasing

Gradient descent




- How to change θ to improve $J(\theta)$?
- Choose a direction in which $J(\theta)$ is decreasing
- Gradient $\frac{\partial J(\theta)}{\partial \theta}$
- Positive \Rightarrow increasing
- Negative \Rightarrow decreasing

Gradient descent in more dimensions

- Gradient vector

$$\nabla J(\underline{\theta}) = \left[\frac{\partial J(\underline{\theta})}{\partial \theta_0} \quad \frac{\partial J(\underline{\theta})}{\partial \theta_1} \quad \dots \right]$$



- Indicates direction of steepest ascent
(negative = steepest descent)

Gradient descent


- Initialization
- Step size
 - Can change as a function of iteration
- Gradient direction
- Stopping condition

Initialize θ

Do {

$$\theta \leftarrow \theta - \alpha \nabla_{\theta} J(\theta)$$

} while ($\alpha \|\nabla J\| > \epsilon$)



Gradient for the SSE

- SSE $J(\underline{\theta}) = \frac{1}{2} \sum_j (y^{(j)} - \underline{\theta} \cdot \underline{x}^{(j)T})^2$

- $\nabla J = ?$

$$J(\underline{\theta}) = \frac{1}{2} \sum_j (y^{(j)} - \theta_0 x_0^{(j)} - \theta_1 x_1^{(j)} - \dots)^2$$

$$\frac{\partial J}{\partial \theta_0} = \frac{\partial}{\partial \theta_0} \frac{1}{2} \sum_j (g(\theta))^2$$

$$= \frac{1}{2} \sum_j \frac{\partial}{\partial \theta_0} (g(\theta))^2$$

$$= \frac{1}{2} \sum_j 2g(\theta) \frac{\partial}{\partial \theta_0} g(\theta)$$

$$\frac{\partial}{\partial \theta_0} g(\theta) = \cancel{\frac{\partial}{\partial \theta_0} y^{(j)}} - \frac{\partial}{\partial \theta_0} \theta_0 x_0^{(j)} - \cancel{\frac{\partial}{\partial \theta_0} \theta_1 x_1^{(j)}} - \dots$$

0 **0**

$$= -x_0^{(j)}$$

Gradient descent

- Initialization
- Step size
 - Can change as a function of iteration
- Gradient direction
- Stopping condition

Initialize θ

Do {

$$\theta \leftarrow \theta - \alpha \nabla_{\theta} J(\theta)$$

} while ($\alpha \|\nabla J\| > \epsilon$)

$$J(\underline{\theta}) = \frac{1}{2} \sum_j (y^{(j)} - \underline{\theta} \cdot \underline{x}^{(j)})^2$$

$$\nabla J(\underline{\theta}) = - \sum_j \underbrace{(y^{(j)} - \underline{\theta} \cdot \underline{x}^{(j)})}_{\text{Error magnitude \& direction for datum } j} \cdot \underbrace{[x_0^{(j)} x_1^{(j)} \dots]}_{\text{Sensitivity to each } \theta_i}$$

**Error magnitude &
direction for datum j**

**Sensitivity to
each θ_i**

Derivative of SSE

$$\nabla J(\underline{\theta}) = - \sum_j (y^{(j)} - \underbrace{\underline{\theta} \cdot \underline{x}^{(j)T}}_{\text{Error magnitude \& direction for datum } j}) \cdot \underbrace{[x_0^{(j)} x_1^{(j)} \dots]}_{\text{Sensitivity to each } \theta_i}$$


- Rewrite using matrix form

$$\begin{aligned} \underline{\theta} &= [\theta_0, \dots, \theta_n] \\ \underline{y} &= [y^{(1)} \dots, y^{(m)}] \end{aligned} \quad \underline{X} = \begin{bmatrix} x_0^{(1)} & \dots & x_n^{(1)} \\ \vdots & \ddots & \vdots \\ x_0^{(m)} & \dots & x_n^{(m)} \end{bmatrix}$$

$$\nabla J(\underline{\theta}) = (\underline{y} - \underline{\theta X}^T) \cdot \underline{X}$$


(Matlab) `>> e = y - th*X'; DJ = e*X; th=th - al*DJ;`

Gradient descent on cost function





Comments on gradient descent

- Very general algorithm
 - we'll see it many times
- Local minima
 - Sensitive to starting point



Comments on gradient descent

- Very general algorithm
 - we'll see it many times
- Local minima
 - Sensitive to starting point
- Step size
 - Too large? Too small? Automatic ways to choose?
 - May want step size to decrease with iteration
 - Common choices:
 - Fixed
 - Linear: $C/(\text{iteration})$
 - Newton's method (we'll return to this...)



SSE Minimum

$$\nabla J(\underline{\theta}) = (\underline{y} - \underline{\theta} \underline{X}^T) \cdot \underline{X} = \underline{0}$$

- Reordering, we have

$$\underline{y} \underline{X} - \underline{\theta} \underline{X}^T \cdot \underline{X} = \underline{0}$$


$$\underline{y} \underline{X} = \underline{\theta} \underline{X}^T \cdot \underline{X}$$

$$\underline{\theta} = \underline{y} \underline{X} (\underline{X}^T \underline{X})^{-1}$$

- $\underline{X} (\underline{X}^T \underline{X})^{-1}$ is called the “pseudo-inverse”

$$\underline{y} \approx \hat{\underline{y}} = \underline{\theta} \underline{X}^T \quad \hat{\underline{\theta}} = \underline{y} \cdot \text{inv}(\underline{X}^T)$$


- If \underline{X}^T is square and independent, these are the same
- If overdetermined, pseudo-inverse gives MSE estimate



Normal equations

$$\nabla J(\underline{\theta}) = 0 \quad \Rightarrow \quad (\underline{y} - \underline{\theta X}^T) \cdot \underline{X} = \underline{0}$$

- Interpretation:
 - $(y - \theta X) = (y - \hat{y})$ is the vector of errors in each example
 - X are the features we have to work with for each example
 - Dot product = 0: orthogonal



$$\underline{y} = [y^{(1)} \dots y^{(m)}]$$
$$\underline{x}_i = [x_i^{(1)} \dots x_i^{(m)}]$$

Matlab SSE

- This is easy to solve in Matlab...



$$\underline{\theta} = \underline{y} \underline{X} (\underline{X}^T \underline{X})^{-1}$$

```
% y = [y1 ... ym]
% X = [x1_0 ... x1_m ; x2_0 ... x2_m ; ...]

% Solution 1: "manual"
th = y * X * inv(X' * X);

% Solution 2: "mrdivide"
th = y / X';           % th*X' = y => th = y/X'
```

More dimensions?



$$\hat{y}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2$$



$$\hat{y}(x) = \underline{\theta} \cdot \underline{x}^T$$

$$\underline{\theta} = [\theta_0 \ \theta_1 \ \theta_2]$$

$$\underline{x} = [1 \ x_1 \ x_2]$$

Nonlinear functions

- What if our hypotheses are not lines?
 - Ex: higher-order polynomials



Nonlinear functions

- Consider the polynomial in x :

$$\hat{y}(x) = \theta_0 + \theta_1 x^1 + \theta_2 x^2 + \theta_3 x^3$$

- This function is still linear in theta
 - Only nonlinear in x ...

- Recall defining $x_0 = 1$

- Let's define $x_p = x^p$

- $x_0 = x^0 = 1$

- $x_1 = x^1 = x$

- $x_2 = x^2$




- ...

$$\hat{y}(x) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3$$

Exactly the same form as before!

Higher-order polynomials

- Fit in the same way
- More “features”







Features


- In general, can use any features we think are useful
- Other information about the problem
 - Sq. footage, location, age, ...
- Polynomial functions
 - Features $[1, x, x^2, x^3, \dots]$
- Other functions
 - $1/x$, $\text{sqrt}(x)$, $x_1 * x_2$, ...
- “Linear regression” = linear in the parameters
 - Features we can make as complex as we want!

Higher-order polynomials

- Are more features better?
- “Nested” hypotheses
 - 2nd order more general than 1st,
 - 3rd order “ “ than 2nd, ...
- Fits the observed data better




Overfitting and complexity





Overfitting and complexity

Simple model: $Y = aX + b + e$




Test data


- After training the model
- Go out and get more data from the world
 - New observations (x,y)
- How well does our model perform?



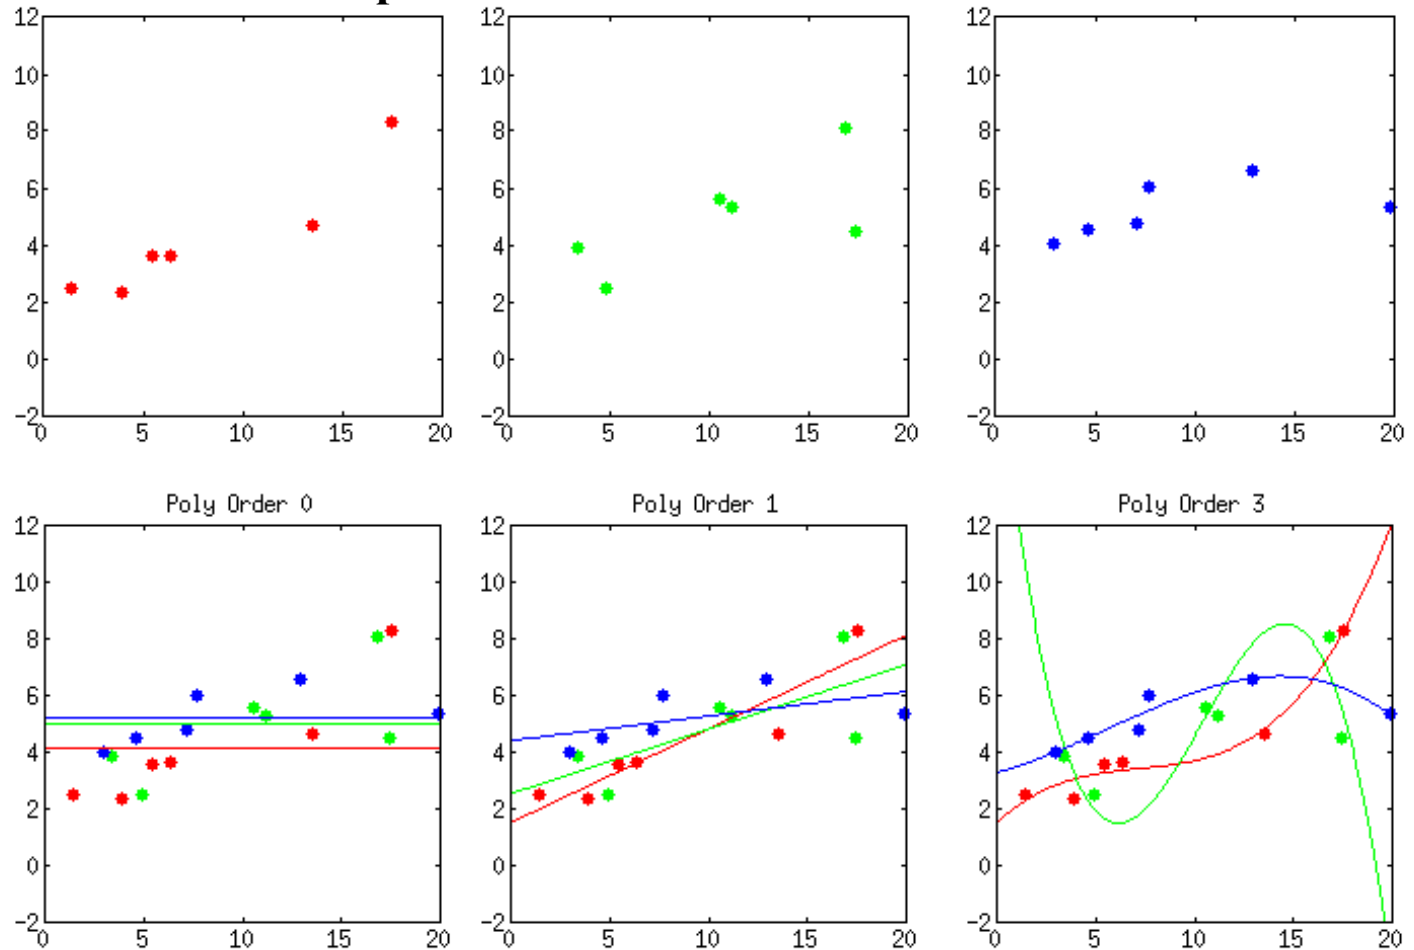
How Overfitting affects Prediction



Bias & variance




Three different possible data sets:



Each would give different predictors for any polynomial degree:

Training versus test error

- Plot SE as a function of model complexity
 - Polynomial order
- Decreases
 - More complex function fits training data better
- What about new data?
- 0th to 1st order
 - Error decreases
 - Underfitting
- Higher order
 - Error increases
 - Overfitting




Detecting overfitting

- Overfitting effect
 - Do better on training data than on future data
 - Need to choose the “right” complexity
- One solution: “Hold-out” data
- Separate our data into two sets
 - Training
 - Test
- Learn only on training data
- Use test data to estimate generalization quality
 - Model selection
- All good competitions use this formulation
 - Often multiple splits: one by judges, then another by you

What to do about under/overfitting?

- Ways to increase complexity?
 - Add features, parameters
 - We'll see more...
- Ways to decrease complexity?
 - Remove features (“feature selection”)
 - “Fail to fully memorize data”
 - Partial training
 - Regularization



Regularization


- Recall $J(\underline{\theta}) = \frac{1}{2}(\underline{y} - \underline{\theta} \underline{X}^T) \cdot (\underline{y} - \underline{\theta} \underline{X}^T)^T$
- Can add “preference” for certain parameters
 - Independent of the data

$$J(\underline{\theta}) = \frac{1}{2}(\underline{y} - \underline{\theta} \underline{X}^T) \cdot (\underline{y} - \underline{\theta} \underline{X}^T)^T + \alpha \underline{\theta} \underline{\theta}^T$$

- New solution (derive the same way)

$$\underline{\theta} = \underline{y} \underline{X} (\underline{X}^T \underline{X} + \alpha I)^{-1}$$




- Problem well-posed for any degree






Regularization

- Compare between unreg. & reg. results

**Alpha =0
(Unreg)**




Alpha =1




Effects of MSE choice

- Sensitivity to outliers



L1 error



Robust cost functions

$$\ell_2 : (y - \hat{y})^2$$

$$\ell_1 : |y - \hat{y}|$$

Something else entirely...

$$c - \log(\exp(-(y - \hat{y})^2) + c)$$

**“Arbitrary” functions can’t be
solved in closed form...
- use gradient descent**

