

Machine Learning and Data Mining

Multi-layer Perceptrons & Neural Networks


Prof. Alexander Ihler
Fall 2012




Linear Classifiers (Perceptrons)

- Linear Classifiers
 - a linear classifier is a mapping which partitions feature space using a linear function (a straight line, or a hyperplane)
 - separates the two classes using a straight line in feature space
 - in 2 dimensions the decision boundary is a straight line

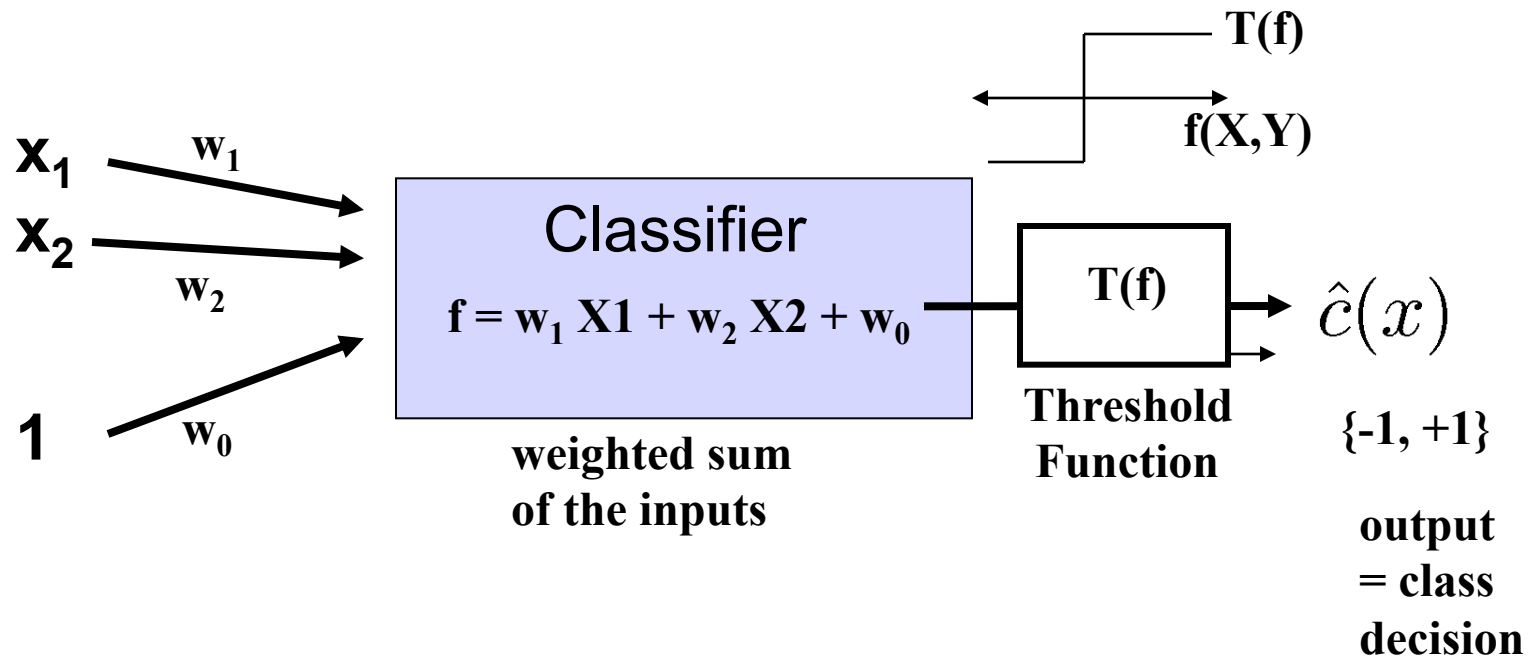
Linearly separable data



Linearly non-separable data



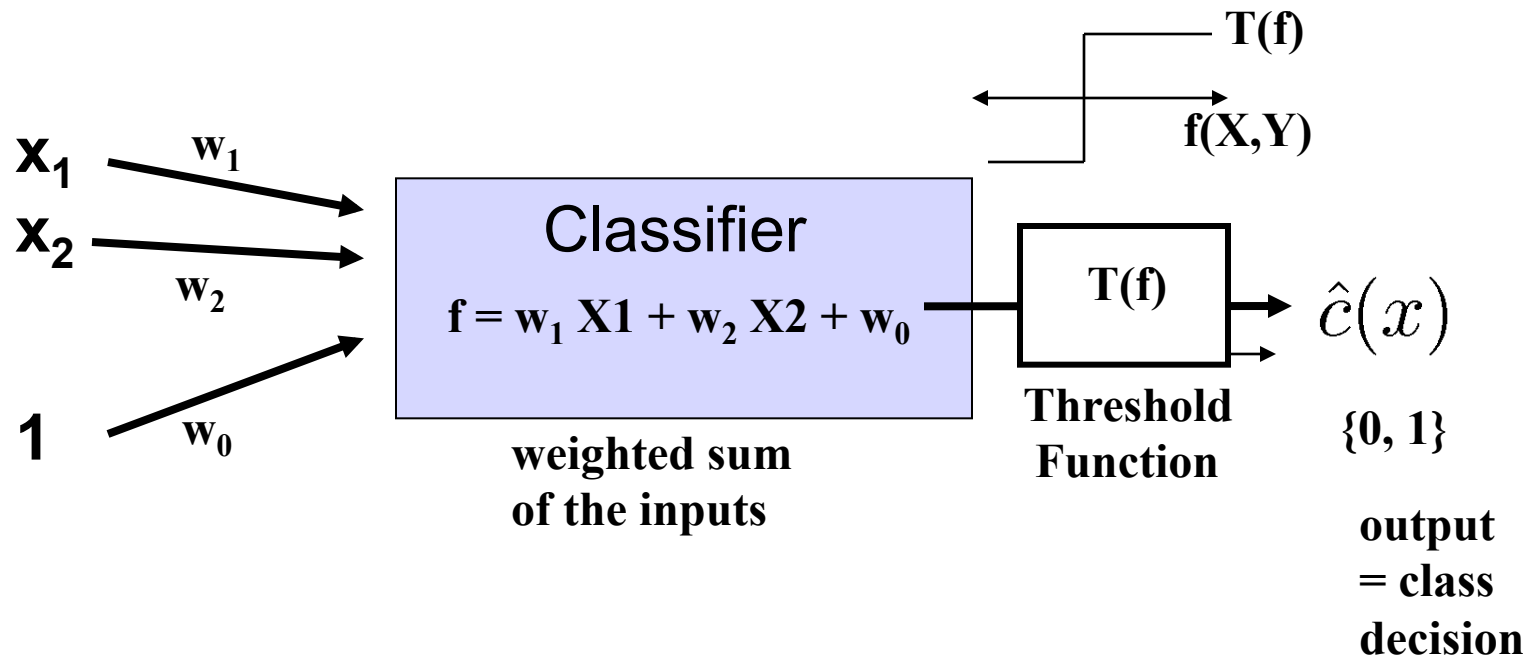
Perceptron Classifier (2 features)




Decision Boundary at $f(x) = 0$

Solve: $X_2 = -w_1/w_2 X_1 - w_0/w_2$ (Line)

Perceptron (Linear classifier)



1D example:





$$\begin{aligned} T(f) &= 0 \text{ if } f < 0 \\ T(f) &= 1 \text{ if } f > 0 \end{aligned}$$

Decision boundary = "x such that $T(w_1 x + w_0)$ transitions"


Features and perceptrons

- Recall the role of features
 - We can create extra features that allow more complex decision boundaries
 - Linear classifiers
 - Features $[1, x]$
 - Decision rule: $T(ax+b) = ax + b > /< 0$
 - Boundary $ax+b=0 \Rightarrow$ point
 - Features $[1, x, x^2]$
 - Decision rule $T(ax^2+bx+c)$
 - Boundary $ax^2+bx+c = 0 = ?$
 - What features can produce this decision rule?



Features and perceptrons

- Recall the role of features
 - We can create extra features that allow more complex decision boundaries
 - For example, polynomial features
- What other kinds of features could we choose?
 - Step functions?



Linear function of features


$$a F1 + b F2 + c F3 + d$$

Ex: $F1 - F2 + F3$



Multi-layer perceptron model

- Step functions are just perceptrons!
 - “Features” are outputs of a perceptron
 - Combination of features output of another





Features of MLPs

- Simple building blocks
 - Each element is just a perceptron f' n
- Can build upwards
 - 2 layer: simple features, complex output
 - 3 layer: complex features
 - 4 layer: even more so...
 - Current research: “deep” hierarchies
- Flexible function approximation
 - Can represent any function arbitrarily closely
 - Given enough hidden units
 - Even a 2-layer with enough hidden nodes

Neural networks


- Another term for MLPs
- Biological motivation
- Neurons
 - “Simple” cells
 - Dendrites sense charge
 - Cell weighs inputs
 - “Fires” axon



“How stuff works: the brain”


Feed-forward networks

- Information flows left – right
- Observed vars input
- Compute hidden nodes
- Compute next layer...
- Info distributed
- Parallel computation
- Alternative: feedback
 - “Recurrent” neural nets
 - Cycles of dependence
 - More complex functions
 - Can have “memory”



Training MLPs

- Observe features “x” with target “y”
 - Push “x” through NN = output is “ \hat{y} ”
 - Error: $(y - \hat{y})^2$
 - How should we update the weights to improve?
-
- Single layer
 - Logistic sigmoid function
 - Smooth, differentiable
 - Optimize using:
 - Batch gradient descent
 - Online gradient descent



Backpropagation

- Just gradient descent...
- Apply the chain rule to the MLP

- Recall: logistic regression

$$\frac{\partial C}{\partial w_i} = -2(y - \hat{y}(w, x)) \sigma'(w, x) x_i$$

- Multi-layer:

— Output layer $\hat{y}_k = \sigma(s_k) = \sigma(w_{k1}^2 h_1 + \dots)$

— Hidden layer $h_j = \sigma(t_j) = \sigma(w_{j1}^1 x_1 + \dots)$


$$\frac{\partial C}{\partial w_{kj}^2} = -2(y_k - \hat{y}_k) \sigma'(s_k) h_j$$

$$\frac{\partial C}{\partial w_{ji}^1} = -2 \sum_k (y_k - \hat{y}_k) \sigma'(s_k) w_{kj} \sigma'(t_j) x_i$$

Logistic regression

$$\begin{aligned} \hat{y}^{(i)} &= \sigma\left(\sum_j w_j x_j^{(i)}\right) \\ &= (1 + \exp(-w x^{(i)}))^{-1} \end{aligned}$$

$$\sigma'(z) = \sigma(z)(1 - \sigma(z))$$



MLPs in practice

- Example: Deep belief nets (Hinton et al. 2007)
 - Handwriting recognition
 - Online demo
 - 10 label \Leftrightarrow 2000 top \Leftrightarrow 500 high \Leftrightarrow 500 mid \Leftrightarrow 784 pixels



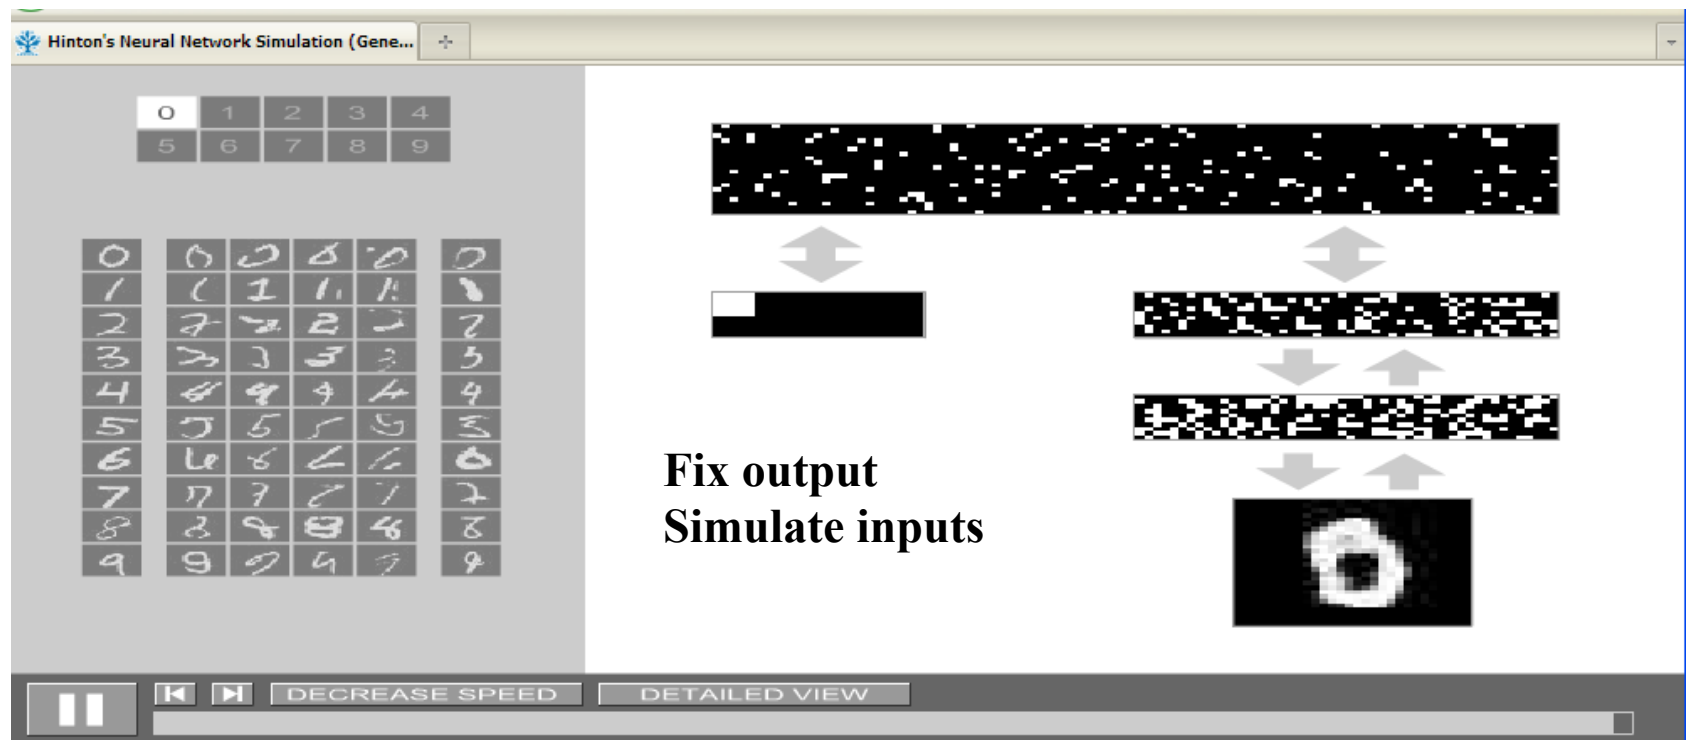
MLPs in practice

- Example: Deep belief nets (Hinton et al. 2007)
 - Handwriting recognition
 - Online demo
 - 10 label \Leftrightarrow 2000 top \Leftrightarrow 500 high \Leftrightarrow 500 mid \Leftrightarrow 784 pixels



MLPs in practice

- Example: Deep belief nets (Hinton et al. 2007)
 - Handwriting recognition
 - Online demo
 - 10 label \Leftrightarrow 2000 top \Leftrightarrow 500 high \Leftrightarrow 500 mid \Leftrightarrow 784 pixels



Neural networks & DBNs

- Want to try them out?
- Matlab “Deep Learning Toolbox”
<https://github.com/rasmusbergpalm/DeepLearnToolbox>
- Also
 - A built-in toolbox for Matlab
 - Have to have a license...
 - Netlab
 - Not updated in some time

Summary

- Neural networks, multi-layer perceptrons
- Cascade of simple perceptrons
 - Each just a linear classifier
 - Hidden units used to create new features
- Together, general function approximators
 - Enough hidden units (features) = any f' n
 - Can create nonlinear classifiers
 - Also used for function approximation, regression, ...
- Training via backprop
 - Gradient descent; logistic; apply chain rule